CPU Virtualization

> Each peach-eater can eat their own *virtual* peach, even though there's
> only one *physical* peach.

The Base Abstraction: Processes

A *process* is a running program. As users, we often want to run multiple processes
at the same time. To make this happen, the operating system must provide the
illusion that there are many CPUs, so that each process can run on one.

This is done through *virtualization* – processes are stopped and started (*context
switching*), sharing the CPU time at the cost of taking longer to run. Operating
systems also have *policies* that determine exactly how decisions like scheduling
are made.

Because a process is by definition a running program, we can understand what
processes are by taking account of everything that happens during execution.

The *machine state* of a process is everything that the program can read of update
while running – all the parts we're interested in when thinking about the process.
Key components of the machine state are:

- *Memory*, which includes the instructions and all the data the program
  touches (its *address space*)
- *Registers*, part of the processor that are read from and written to over the
  course of program execution (including the *program counter*, *stack pointer*,
  and *frame pointer*)
- Various IO devices, like persistent storage

The Process API

In any implementation, the OS needs to provide a basic interface for interacting
with processes, supporting operations:

- *Create*
- *Destroy*
- *Status*
- *Wait*, which can allow one process to wait for another to terminate
- Miscellaneous control, including support for *suspending* and *resuming*
  processes

Process Creation

Process creation involves taking bytes on disk that specify instructions and
turning them into a running program. First, data is *loaded* into memory,

including all of its static data. On modern operating systems, this is done *lazily*. Next, memory is allocated for the *stack*, possibly initialized with values (eg. C's `argc` and `argv`). *Heap* data might also be allocated, and depending on the OS there might be additional initialization tasks (like opening file descriptors). Finally, execution begins from `main`, and control is transferred to the CPU to begin execution.

Process States

Processes have three high-level states:

- Running – executing instructions
- Ready – ready to run but not currently being executed
- Blocked – not ready to run until some other event takes place (for example, waiting for an IO request to resolve)

Transitions between ready and running (*scheduled* vs *descheduled*) happen at the discretion of the OS based on decisions of the *scheduler*.

Data Structures

A few key data structures are used to keep track of state in typical operating systems. For example, a *process list* tracks process state. This includes the *register context* (register contents) of each stopped process, process states, etc.

Aside: The Unix Process API

The core of the Unix process API is three functions: `fork()`, `exec()`, and `wait()`.

Fork

`fork()` can be called by a process to create a clone (exact copy) of the calling process. The function returns twice, with value `0` in the child process and a non-zero positive process id of the child in the calling parent. This process is *nondeterministic* – the parent might run and terminate before the child, or vice-versa.

Wait

`wait()` can be used to have the parent block until the child finishes executing.

Exec

`exec()` is used to run a program that is different from the calling program (which we want to do basically always). The call accepts as arguments the name of an executable (like `wc`) and some arguments, and then *transforms* the currently running process into the one specified by the arguments. Thus, successful calls to `exec` never return, so code after `exec` isn't run.

Motivation

Motivation for the API can be understood by considering how it is more useful than alternatives in practice. For example, having distinct `fork()` and `exec()` calls makes it possible for a shell to modify the process environment before running a program. For example, shell redirection might involve setting up some file descriptors before transferring control to the running process.

Process Control and Users

Beyond the main three functions, Unix also provides interfaces for interacting with processes. For example, `kill()` is used to send *signals* to processes including `SIGINT` (interrupt) which generally terminates process and `SIGSTOP` which typically suspends them. Processes can also "catch" signals with `signal()`, allowing arbitrary code to be run in response to them.

To remain sound alongside core OS goals like security and isolation, signals can't be sent arbitrarily. If this were the case, a rogue process could terminate all others on the system. Instead, the OS maintains a notion of a user, who must log in, and can then send signals to processes they have started, but not those started by other users on the machine.

Mechanism: Limited Direct Execution

CPU virtualization is achieved through *time sharing* – a process is run for a bit, then swapped out, etc. The OS must do this while maintaining high *performance* and full *control*.

Limited Direct Execution

To avoid paying a high performance penalty, user programs are run directly on the CPU. However, this poses two challenges. First, the OS must prevent the process from doing bad things. Second, in order to implement *time sharing* the OS must have some way to take back control.

Problem 1: Restricted Operations

If operations weren't restricted, a user process could do all sorts of malicious things like reading or writing anywhere on disk. To prevent this, a new processor mode called *user mode* is introduced which isn't allowed to do things like make IO requests without raising an exception (as distinct from *kernel mode* which is unrestricted).

When user processes need to perform privileged operations, they issue a *system call.* With this system call, the user includes a system call number (among the values supported by the OS). In response, the OS executes a *trap* instruction, jumping into the kernel and raising the privilege level and conditionally executing the call. The kernel refers to a *trap table* to tell which code to execute (i.e. which *trap handler*) based on the exception that was raised. This table is initialized with the OS boots up and remains constant until the program is rebooted.

After, a *return from trap* instruction is executed, reducing the privilege level and returning control to the user process. Logistically, this involves saving the caller's registers so that the program can return to normal execution after completion.

Restricted operations are thus supported by setup at boot-time along with having two privilege levels.

Problem 2: Switching Between Processes

Because the OS cedes control in the direct execution model, it isn't obvious how it will be able to take it back. Early systems attempted a *cooperative approach*, where the OS trusts processes to give up control. This obviously is problematic because it can't deal with malicious or problematic programs that don't cede control.

Again, hardware support is used in the form of a *timer interrupt.* At configured (at boot-time) intervals, a timer will go off, transferring control to an interrupt handler and thus granting control to the OS.

After a timer interrupt occurs, the OS *scheduler* decides whether to cede control back to that running process to switch to a new one. *Context switching* requires saving registers and restoring those from the newly-running process.

All this is complicated by the fact that an interrupt could occur during the processing of a system call, for example. Solutions to this sort of problem might involve locking, or temporarily disabling interrupts.

Scheduling: Introduction

Policies govern how operating systems make scheduling decisions.

Workload Assumptions

The *workload* is the set of all processes running on a system. Schedulers must make some assumptions about their workloads, and also employ a *scheduling metric* to compare different possible policies. For example, a simple metric is turnaround time, the amount of time a job takes to complete. Alternatively, we could assess *fairness* (or some combination).

Simplest: FIFO

The simplest possible policy is FIFO, which serves processes in the order they arrive. However, this is clearly not optimal because serving a single long-running process while a number of short processes wait can drive up average turnaround time (know as the *convoy effect*).

Shortest Job First (SJF)

Just running the shortest jobs first significantly reduces the average turnaround time (to an optimal level, if we could assume that all jobs arrive at the same time). However, SJF falls apart if jobs arrive at different times – the longest job could arrive first and we'd be back where we started.

Shortest Time-to-Completion First (STCF)

We can deal with this if the scheduler is given the ability to preempt jobs. When the shorter jobs arrive after the long one has started, if running them right away would let them finish earlier, then the scheduler could preempt the long running job and only finish it after processing the short ones. This again improves turnaround time relative to SJF

Another Metric: Response Time

Interactive systems demand that we pay attention to response time, motivating a new metric. By this metric, it's important that a process at least be started soon after its arrival. The previously processed systems are adequate for batch-processed systems, but not so much for interactive systems – we don't want cases where one process has to wait for two others to finish completely before it gets started at all.

Round Robin

Motivated by improving average response time, Round Robin runs jobs for *time slices* (some multiple of the timer-interrupt period) and then switches to the next, cycling through a run queue until jobs are finished. By cycling through

each job quickly, rather than running one after another completion, RR achieves much improved average response time. Of course, we improve by this metric the shorter we make the *time slices*, but must avoid making them too small where the cost of context-switching would dominate.

However, Round Robin is a poor policy by the turnaround time metric. Intuitively, this makes sense – each job gets stretched out as long as it can. This makes round robin nearly the worst possible policy by this metric. However, it is *fair* – CPU resources are evenly distributed among active processes.

Considering IO

IO complicates things, because a single process ends up having periods where no useful work can be done. To get around this, the CPU treats each "burst" (normal period, not waiting on an IO request) as a separate job. This allows for *overlap*, which is generally good for resource utilization.

Considering Unknown Time Requirements

Much of the prior discussion makes sense under the unrealistic assumption that we can know how long each job will take to run. Of course, in practice this isn't possible. It will turn out that operating systems use information from the recent past to make predictions about the future, here (known as a *multi-level feedback queue*).

Scheduling: The Multi-Level Feedback Queue

The aim of MLFQ is to optimize turnaround while simultaneously minimizing turnaround time, all without knowledge like how long a process will take to run.

MLFQ: Basic Rules

An MLFQ consists of a number of queues, each with different priority levels. Jobs that are ready to be run on the queue, and at any given moment the OS chooses from a process from a high priority queue to run. Among jobs present in the same queue (i.e. they have the same priority), RR is used.

Job priority is *varied* based on *observed behavior*, and thus *changes* over time.

Changing Priority: Attempt 1

Consider the rules

- New jobs are given top priority

- Jobs that take up an entire time slice while running are demoted
- Jobs that don't stay where they are

Intuitively, what happens is the system first assumes that the process is short-running, and continually updates its assumption and correspondingly demotes it (so that new short jobs that come in have a chance to finish) as it realizes that how long it's going to take.

This system is obviously good for short-running jobs, but it's also good for jobs that do a lot of IO, which look to the scheduler like a stream of short-running jobs. However, loads of interactive jobs can prevent long-running processes from ever running (cause them to *starve*). A malicious application developer could easily game this to monopolize the CPU, relinquishing control just before the end of each time slice to retain high priority. Also, a process that changes its behavior can't be re-promoted.

The Priority Boost: Attempt 2

To combat starvation, we can try to periodically boost the priority of all jobs, like by promoting everything to the top queue periodically. At minimum, this guarantees against starvation, since a long running process will run at least a bit after each boost. Setting the constant is difficult, since we need to balance between not letting long-running processes starve and allowing interactive jobs sufficient access to the CPU.

Better Accounting: Attempt 3

One way to prevent CPU gaming mentioned in Attempt 1 is to track an allotment of CPU time. Thus, relinquishing the CPU just before the timer would no longer be useful, as the process would be demoted soon afterwards.

Tuning MLFQ

There are also questions about how to set other MLFQ parameters. For example, we need to decide how many queues there should be, and how long time slices should be for each queue. For example, many implementations have short time slices for high-priority queues, but longer ones for the long-running processes that inhabit lower priority ones.

Scheduling: Proportional Share

Proportional share schedulers optimize for fairness (guaranteeing each job a certain percentage of CPU time) rather than turnround or response time. This is known as *lottery scheduling*: a lottery is periodically held to determine which process gets to run next.

Tickets Represent Shares

Tickets represent the share of a resource that a process should receive, so if process A has 75 tickets and B has 25, then A gets 75% of CPU time probabilistically.

Ticket Mechanisms

Tickets can be manipulated. For example, a user can allocate their tickets to their jobs however they choose (*ticket currency*). Tickets can also be *transferred*, and even their value can be *inflated* within a group of trusted processes.

Implementation

Lottery scheduling can be implemented extremely easily. All we need is a random number generator, knowledge of the total number of tickets, and access to the process list. Then, finding the winner is just like indexing into a list.

Unfairness

In expectation, a random scheduler will allocate resources in accordance with the ticket distribution. Because this doesn't always happen, especially with small sample sizes, we can quantify the drift using an *unfairness metric*.

Ticket Assignment

It isn't obvious how to allocate tickets, particularly without getting information from the user.

Stride Scheduling

To deal with the unfairness that can arise due to randomness, there exist deterministic algorithms like *stride*. However, the tradeoff these make is that they require that we maintain some global state.

Linux's Completely Fair Scheduler

Linux's scheduler aims to evenly divide CPU among competing processes. To achieve this, is uses a concept called `vruntime`. Running processes accumulate `vruntime`, and whenever a scheduling decision is made, the process with the *lowest* `vruntime` is chosen. Parameters tune things like how frequently the scheduler preempts running processes, with the tradeoff being between perfect

fairness and performance. Other parameters like `min_granularity` prevent the scheduler from preempting jobs too frequently, especially when there are many competing jobs at a time.

Rather than using tickets, Unix has a notion of "niceness", where programs that have more negative niceness have higher priority. Red-black trees (balanced binary trees) are used to ensure that the scheduler must do minimal work while selecting the next process to run.

Multiprocessor Scheduling

Multicore systems face classes of problems that don't exist on single core ones. For example, we can consider the problem of *cache coherence*. On a single core system, fetching a value from cache is a straightforward process. However, with multiple processors and multiple caches, we can easily see that caches can drift out of correctness. Hardware support solves this basic problem via *bus snooping*: each cache pays attention to all updates, and invalidates its own entries when needed.

There's also the need to use *locks* (or *lock-free* data structures) to avoid data synchronization issues. Lastly, multiprocessor schedulers must also consider *cache affinity*: it's better to run a process on the same CPU each time if possible, where the required data might be ready in cache.

Easy Solution: Single-Queue Scheduling (SQMS)

Basically the easiest possible solution to the problem of multicore scheduling is using a single queue (SQMS). However, this solution is not *scalable*. For example, locks are required to ensure that operations like access to the single queue run properly, resulting in a performance cost. Also, SQMS completely ignores the problem of cache affinity.

Multi-Queue Scheduling (MQMS)

Using multiple queues solves many of the issues of SQMS. New jobs are placed onto one scheduling queue according to some policy, and remain there. There is less or no need for locks, and thus also less cache contention. Moreover, processes generally stay in one queue, which leads to good cache affinity.

However, we introduce a new problem of *load balancing*. In the worst case, for example, we could have all the jobs on one of two CPUs finish, leading one CPU to sit idle while the other continues to work (poor utilization). The solution is *migration*, which moves jobs from one CPU to another to achieve load balance. This is implemented with *work stealing*, where queues peek at others and take jobs from them if they are significantly less full. Naturally, there's a tradeoff

between the overhead of checking queues frequently and not checking frequently
enough.